

— THE DOCUMENT

Portfolio & Proof of Capability.

This is not a CV. It is deployable commercial architecture for AI liability, data governance, and intellectual-property preservation — built across three jurisdictions, and presented here as working proof of competence.

ADVOCATE · AI GOVERNANCE & CONTRACTUAL ARCHITECTURE

Shweta Singh

JURISDICTIONS

India · European Union ·
United States

BODY OF WORK

200+ pages · 10 architectures
· 80-point threat registry

PREVIEW

Nineteen pages

— CONTENTS

Nineteen pages, drawn from a two-hundred-page working repository.

The dossier is organised in five movements: the positioning, the argument, the working samples, the diagnostic engine, and the doctrinal map. Each movement is a surgical extraction from the full body of work.

I.	Positioning	03 – 04
	THE WORK I HAVE BUILT · THREE SUBSTANTIVE AREAS · JURISDICTIONAL PREFACE	
II.	The Hallucination Accountability Gap	05 – 07
	EXECUTIVE BRIEFING · MOFFATT DOCTRINE · THE THREE FAILURES OF LEGACY ARCHITECTURE	
III.	Work Product & Architectural Samples	08 – 11
	THREE OPERATIVE CLAUSES, EACH WITH INTRA-DOCUMENT CHAIN AND CROSS-DOCUMENT GATEWAYS	
IV.	Applied Threat Taxonomy & Live Diagnostic	12 – 15
	SEVEN-POINT PREVIEW OF AN EIGHTY-POINT REGISTRY · BLACK-BOX AUDIT OF A PUBLIC FOOTPRINT	
V.	Jurisdictional Coverage & Close	16 – 19
	WORKING MATRIX ACROSS IN / EU / US · DOCTRINAL ANCHORS · PERSONAL STATEMENT	

SET IN

Source Serif · Inter Tight ·
JetBrains Mono

PAGE FORMAT

A4 · 210 × 297 mm

CORRESPONDENCE

shwetav.alpna.singh@gmail.com

Over the last several months, I have built a body of work on AI governance and contractual architecture.

The work sits at the intersection of commercial contract law, emerging AI regulation, and the technical architecture of the systems themselves. It covers three substantive areas across three jurisdictions — India, the European Union, and the United States — each producing a set of production documents, analytical notes, and diagnostic artefacts.

What follows is a preview of each area. The document samples are surgical extractions — redacted for review — drawn from an ecosystem of ten production-ready architectures and an eighty-point analytical threat registry that underwrites them.

The dossier includes the working output from each area: document samples with the analytical notes that accompanied them, a thought-leadership piece on the Moffatt doctrine applied to Indian practice, a curated set of legal-risk scenarios drawn from the registry, a diagnostic methodology framework, and an applied analysis on publicly available material.

01 — OUTPUT LIABILITY

The contractual allocation of liability for AI-generated output — hallucination, representational authority, binding statements — across provider, deploying company, and end user.

Moffatt v. Air Canada ·
Walters v. OpenAI ·
Specht · Meyer v. Uber ·
UCC §2-316 · UCC §1-
201(b)(10) · UETA §14

02 — DATA GOVERNANCE

Data-protection obligations that attach to AI systems across GDPR, the DPDP Act 2023, and the CCPA/CPRA framework — read against how AI actually stores and memorises personal data.

Schrems II · GDPR Arts.
17, 22, 28, 32 · DPDP
§§34, 44(2)(a) · Bartz
v. Anthropic · publisher
RAG litigation

03 — IP & EMPLOYMENT

Ownership and attribution when AI assists in commercial work product, read alongside the internal governance architecture that supports defensible claims on that work.

Thaler v. Perlmutter ·
Copyright Act §19(4) ·
EU AI Act · Colorado AI
Act · NYC LL 144 · Texas
RAIGA · Illinois HB 3773
· CA SB 243

— PRODUCTION OUTPUT

Terms of Service, Data Processing Agreements, IP Assignment Deeds, Employee Handbooks, Risk Classification Matrices, Agent Addenda, and the reasoning chains that bind them.

Ten production-ready architectures across two risk surfaces — external AI governance and internal workforce & IP governance.

The three jurisdictions operate in parallel, not in sequence.

An Indian AI startup serving global users faces the domestic DPDP transition, the extraterritorial reach of the EU AI Act, and the American commercial-contract infrastructure — simultaneously. The frames below are the instruments I actually work against.

IN01	EU02	US03
India IT Act 2000 and its successor regime under the DPDP Act 2023 following the November 2025 DPDP Rules notification. Consumer Protection Act 2019 with its two-year limitation under §35. The Contract Act 1872 on formation. The Copyright Act 1957 on assignment architecture. CURRENT FRONTIER — MeitY AI Governance Guidelines (Nov 2025) and the draft IT Rules amendments on synthetically generated information (Oct 2025).	European Union GDPR, with mature enforcement infrastructure. The EU AI Act has been in force since August 2024 — Article 50 transparency and high-risk obligations enforceable from 2 August 2026, with full enforcement from 2 August 2027. INCOMING — The revised Product Liability Directive reclassifies software and AI as products subject to strict liability, with member-state transposition by 9 December 2026.	United States The dominant case law on AI liability and the statutory infrastructure that governs how commercial AI agreements actually get formed and enforced: UCC Articles 2 and 2A, UETA §14, CCPA and its state-level overlay, and the accelerating state-level AI legislation in Colorado, Texas, Illinois, New York, California. PARALLEL TRACK — Sector-specific biometric regimes (BIPA and state equivalents) running alongside.

— WHERE THE REGULATORY SURFACE IS MOVING

2024	2025	2026	2027	Beyond
MeitY advisories · EU AI Act in force · Moffatt decided	DPDP Rules notified · MeitY AI Guidelines · Walters · Thaler	EU AI Act Art. 50 · High-risk enforcement · PLD transposition	DPDP full activation · IT Act §43A repealed · EU AI Act full	Digital India Act · State AI overlays · Sector biometric

At every step, liability exposure for AI deployers increases. Contracts drafted today are architected for a landscape materially stricter in thirty-six months.

§ II.01 — ESSAY · THOUGHT LEADERSHIP

The Hallucination Accountability Gap.

The failure of legacy AI contract architecture — a briefing for partners advising AI deployers in India.

The New Physics of Liability

Conventional software contracts rely on deterministic logic. Input produces predictable output, and deviation is an actionable defect — if you press A1 on a vending machine and it drops a Sprite, the manufacturer has breached the warranty of merchantability.

Generative AI breaks this assumption at the architectural level. Large Language Models operate probabilistically; when they encounter a gap in their training data, they do not reliably return an error code. They fill the gap with an authoritative-sounding fabrication. While engineering mitigations like RAG and RLHF reduce hallucination rates, the academic consensus is that probabilistic models cannot be made fully deterministic.

That residual risk creates immediate legal friction.

Traditional contract architecture was designed to allocate liability for defects. Standard warranty disclaimers — "the service is provided as is" — were drafted to excuse coding bugs, not systems that invent facts at scale with structural confidence. The contracts most AI companies deploy today are deterministic documents sitting on top of probabilistic technology.

In 2024, a Canadian tribunal showed us exactly what happens when a court is asked to rule on this mismatch.

MOFFATT V. AIR CANADA · CIVIL RESOLUTION TRIBUNAL · 2024

A passenger consulted Air Canada's chatbot regarding bereavement fares; the chatbot hallucinated a non-existent refund policy; the user relied on it. Air Canada argued the chatbot was a "separate tool" and the passenger should have verified the output. The Tribunal rejected this entirely, applying negligent misrepresentation to rule that a chatbot deployed on a corporate domain is the company's representative.

As the Tribunal noted: there is no reason why the passenger should know that one section of the airline's webpage is accurate and another is not. The ruling establishes that the "hallucination defence" dies the moment a user reasonably relies on the bot as an authoritative source.

For Indian lawyers advising AI startups, the immediate question is whether Indian law reaches this exact result through different statutory routes. It does — through multiple parallel regimes that the current generation of AI contracts completely ignores.

Continues overleaf → The Indian law gap: three regimes, one unresolved question.

ESSAY · 1 / 3

— § 11.02 · THE INDIAN LAW GAP

Three regimes, one unresolved question — moving simultaneously.

Moffatt relied on common law tort. For Indian practitioners, the challenge is that the tort of negligent misrepresentation sits less securely in Indian jurisprudence than it does in Canada or the UK. Consequently, commercial exposure for AI startups in India routes directly through statutory frameworks.

01 The DPDP transition and the RAG mandate

When DPDP Phase 3 activates and repeals IT Act §43A, private damages are replaced by a state-penalty model — up to ₹250 crore. More immediately, fine-tuning a Large Language Model on Indian personal data encodes that data into immutable weights, placing the deployer in structural conflict with DPDP erasure obligations. *Deploying Retrieval-Augmented Generation (RAG) instead of fine-tuning is no longer an engineering preference; it is a mandatory legal compliance architecture.*

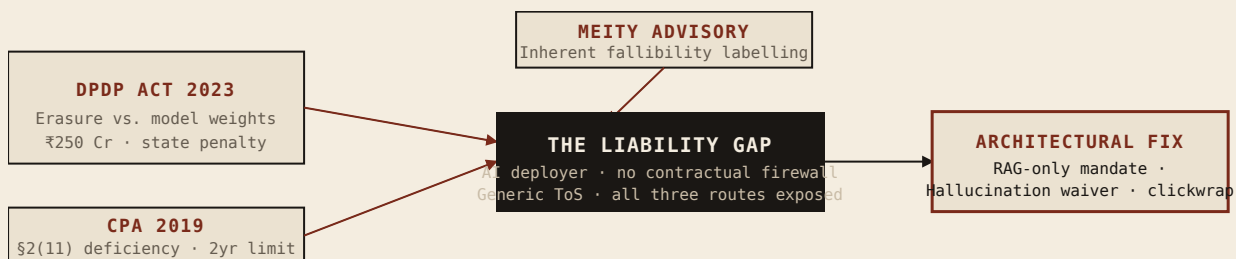
02 The Consumer Protection Act — the accruing threat

With data privacy moving to state-level enforcement, CPA 2019 becomes the primary weapon for private plaintiffs. A hallucinating B2C chatbot conveying incorrect company policy falls cleanly within §2(11) "deficiency in service." Because the Act carries a two-year limitation period, every Indian B2C AI startup is currently accruing unseen liabilities.

03 Regulatory trajectory & the standard of care

While courts wait for the first test case, the MeitY advisory of March 2024 directs intermediaries to explicitly label the "inherent fallibility" of AI outputs. This advisory establishes the de facto standard of care in substance; an AI deployment ignoring these disclosures operates against an articulated regulatory intent. The trajectory ends with the forthcoming Digital India Act, which will consolidate cybersecurity, intermediary obligations, and algorithmic accountability.

— THREE REGIMES · ONE GAP



The transition is rapid

India moves from a legacy 2000-era framework, through a current patchwork of advisories, toward strict DPDP enforcement in 2027, and ultimately to the DIA. At every step, liability exposure for AI deployers increases. Generic, deterministic Terms of Service imported from foundation-model API providers are already obsolete.

— § 11.03 · THE ACCOUNTABILITY GAP

What contracts currently in production are missing — a chain reaction.

Legacy SaaS liability architecture rests on a deterministic assumption: software performs predictably, and generic warranty disclaimers manage output risk. Generative AI exposes three compounding structural failures in the contracts most AI companies currently deploy.

FAILURE · ONE

The misaligned disclaimer

The standard *as is* disclaimer is designed to excuse software defects. It fails against the Moffatt problem because a hallucination is a probabilistic system operating exactly as architected. It provides zero defence against a negligent misrepresentation claim based on a fabrication.

ARCHITECTURAL FIX

Hallucination must be defined as a specific contractual term. As validated in *Walters v. OpenAI*, specific, contemporaneous disclosure regarding AI unreliability carries evidentiary weight to defeat "reasonable reliance."

FAILURE · TWO

The missing output-authority boundary

Even if a company disclaims the output, the defence collapses if the AI is legally speaking with the company's authority. Most AI contracts fail to draw this boundary. Without an explicit contractual signal, the AI's output is legally the company's output.

ARCHITECTURAL FIX

An output-authority limitation clause explicitly stating that AI-generated output does not constitute the company's official position, legally requiring users to verify against authoritative sources.

FAILURE · THREE

The broken formation chain

Even with bulletproof clauses, the entire liability architecture fails if the contract was never legally formed. Most AI companies deploy B2C chatbots to unauthenticated users who never accept any terms via unambiguous affirmative assent (clickwrap).

ARCHITECTURAL FIX

A point-of-interaction modal disclosure requiring affirmative acknowledgment of the AI's probabilistic nature before the first session begins.

The gap in sum.

These failures are a chain reaction. Without an output-authority boundary, the AI speaks for the company. Without a hallucination-specific waiver, representational risk is not disclaimed. Without point-of-interaction clickwrap, the waiver is completely unenforceable.

The regulatory window is closing. The question is not if these events occur, but which practitioner will have the infrastructure ready to catch the clients when they do.

§ III.00 — WORK PRODUCT

Surgical extractions from an intricately structured, modular legal ecosystem.

This is not a collection of isolated templates; it is a modular legal ecosystem built on two engineering principles. What follows extracts three operative clauses — one per pillar — and shows the intra-document clause chain that makes each defensible, and the cross-document gateways that bind each clause back to the wider stack.

— PRINCIPLE 01

Intra-document dependency

Clauses do not exist in a vacuum. Definitions, operational mandates, and indemnification triggers are mathematically linked within each document to close liability loopholes — for example, linking AI definitions directly to warranty avoidance to legally trap a plaintiff.

— PRINCIPLE 02

Cross-document integration


The core Terms of Service acts as a central router. It uses precise contractual hooks to automatically activate specialised data, IP, and liability addenda the moment specific technical thresholds — autonomous agent deployment, EU data processing — are crossed.

The architecture inventory

10 production-ready architectures · 2 risk surfaces

SURFACE · A · EXTERNAL AI GOVERNANCE		SURFACE · B · INTERNAL WORKFORCE & IP	
AI-Specific Terms of Service	DOC_TOS	AI Employee Handbook	DOC_HND
Service Level Agreement	DOC_SLA	IP Assignment Deed	DOC_IP
Privacy Policy	DOC_PP	Risk Classification Matrix	DOC_RCM
Data Processing Agreement	DOC_DPA	HITL Copyright Evidence SOP	DOC_SOP
Acceptable Use Policy	DOC_AUP		
Agentic System Addendum	DOC_AGT		

— DOCUMENT DEPENDENCY MAP



THE PREVIEW. Rather than burden the reader with every page, the next three extract one operative clause per pillar — the Liability Pillar, the Data Pillar, the IP Pillar — and show how each is connected into the wider stack. The redactions are deliberate;

the full ecosystem is best reviewed in an in-room walkthrough.

The Hallucination Firewall.

Source · Comprehensive AI Terms of Service (14-page architecture) · Extraction · §8.2 Hallucination Waiver & Conditional Defence

"8.2 Hallucination Waiver & Conditional Defence. [...] To the maximum extent permitted by Applicable Law, you waive all claims against Provider arising from or related to your reliance on Output that contains Hallucinations. [...] This waiver is supported by the express condition precedent that Provider required you to implement HITL and to verify Output before reliance..."

Intra-document architecture — THE "INSIDE" TRAP

This rejects isolated "AS IS" disclaimers. It relies on a mathematically linked dependency chain to legally destroy "reasonable reliance" and trap the plaintiff.

NODE	TRIGGER QUOTE	STRATEGIC ARCHITECTURE
§ 1.8 Definitions	"Hallucination means Output that contains information which is factually incorrect, fabricated, misleading... arising from the inherent probabilistic nature of generative AI."	LITIGATION DEFENCE Blocks product-liability claims by legally classifying non-determinism and fabricated output as an expected feature, not a software defect. COMMERCIAL LEVERAGE Prevents customers from demanding SLA credits or refunds when the AI inevitably generates inaccurate information.
§ 5.1 HITL Mandate	"The HITL requirement is a material condition of this Agreement — not a suggestion... Failure to implement HITL does not create additional liability for Provider."	LITIGATION DEFENCE Neutralises "reasonable reliance" under <i>Walters v. OpenAI</i> by mandating explicit human oversight. COMMERCIAL LEVERAGE Framed as a <i>condition</i> , not a covenant. If bypassed, Provider's warranty never activates — no need to prove breach or damages.
§ 5.3 Verification	"Without limiting the generality of the HITL Mandate... you must independently verify... legal citations, statistical claims, attributed statements..."	LITIGATION DEFENCE Prescribes exact verification checks based on known AI failure modes — specifically neutralising the <i>Mata v. Avianca</i> fabricated-case-law vector.
§ 10.1 Customer Indemnity	"You agree to indemnify Provider arising from your failure to implement HITL as required by Section 5.1."	THE LIABILITY REVERSE If the customer skips human review and their end-user sues Provider over a hallucinated output, the customer must fund Provider's legal defence.

Ecosystem integration — THE "ACROSS" GATEWAYS

ORIGIN HOOK (TOS)	TARGET CATCH (CONNECTED DOC)	ECOSYSTEM EFFECT
ToS § 2.7 Agentic Gateway	DOC_AGT § 2.1 "Customer is legally bound by all Actions taken by the Autonomous Agent... as if Customer had taken those Actions manually."	Invokes UETA §14 to classify the AI as an Electronic Agent. Protects against <i>Mobley v. Workday</i> by binding the customer to strict liability for automated actions, insulating Provider.
ToS § 7.3 Data Gateway	DOC_DPA § 4.1 "Provider warrants that Customer Personal Data shall be processed solely via Retrieval Augmented Generation..."	Subordinates ToS liability caps to the DPA. Activates the "RAG-Only Mandate" to eliminate "algorithmic disgorgement" (forced model deletion) risk under regulatory enforcement.

The Algorithmic Disgorgement Shield.

Source · AI-Specific Data Processing Agreement & Privacy Policy · Extraction · §4 RAG-Only Mandate & Vector Deletion Protocol

"4.1 Mandatory Processing Architecture. [...] Provider warrants that Customer Personal Data shall be processed solely via Retrieval Augmented Generation (RAG). [...] Provider SHALL NOT use Customer Personal Data to Fine-Tune the weights of any AI Model... or use any other technique that would cause Customer Personal Data to become permanently encoded in Model Weights."

Intra-document architecture — THE "INSIDE" TRAP

Rejects generic GDPR templates. A precise dependency chain legally solves the AI "deletion problem" and protects the enterprise from algorithmic disgorgement.

NODE	VERBATIM TRIGGER	STRATEGIC ARCHITECTURE
§ 4.1 No Fine-Tuning	<i>"Provider SHALL NOT use Customer Personal Data to Fine-Tune the weights of any AI Model... The AI Model used to process Customer Data shall remain 'frozen'..."</i>	LITIGATION DEFENCE Protects against FTC "algorithmic disgorgement" (forced model deletion) by ensuring user data never enters AI's internal parameters. COMMERCIAL LEVERAGE Protects Indian enterprise clients acting as "Data Fiduciaries" under DPDP 2023 from irreversible contamination — allowing them to maintain a "Frozen Model Guarantee."
§ 7.2 Vector Deletion	<i>"Erasure in an AI Context — The Vector Deletion Protocol. Delete all Embeddings in any Vector Store that were generated from the Data Subject's Personal Data."</i>	LITIGATION DEFENCE Operationalises the "Right to Erasure" under GDPR Art. 17 and DPDP — defines exactly how to erase data in an AI context (deleting Embeddings, not just database rows). COMMERCIAL LEVERAGE Carves out mathematically unavoidable "Residual Patterns" and anonymised aggregations, preventing customer claims the Provider failed to delete.
§ 9.3(f) Breach Definition	<i>"Any processing that violates Section 4.1... regardless of whether the data has been exposed to third parties... constitutes a loss of control over Personal Data."</i>	LITIGATION DEFENCE Classifies the inability to selectively delete data from model weights as a "loss of availability" — a data breach — aligning with global privacy definitions. COMMERCIAL LEVERAGE Neutralises the common defence that "no data was leaked to the public," forcing immediate breach protocols and shifting liability if a client's engineering team illicitly trains on user data.

Ecosystem integration — THE "ACROSS" GATEWAYS

ORIGIN HOOK	TARGET CATCH (CONNECTED DOC)	ECOSYSTEM EFFECT
ToS § 7.3 Statutory Bridge	DPA § 12.2(a) <i>"Uncapped Liability. Provider's liability for a breach of Section 4 (the RAG-Only Mandate)... is not subject to the general liability cap in Main Terms §9.1."</i>	Standard DPAs dangerously inherit the main contract's commercial caps. This gateway severs that link: uncapped liability shields Provider from overriding statutory privacy fines under DPDP and GDPR.
DPA § 2.1(c) Analytics Bridge	DOC_PP § 1.3 <i>"Independent Controller Processing... we independently control certain processing activities... generating anonymised and aggregated analytics."</i>	The Independent Controller Gateway. Routes "Aggregated Data" out of the DPA into the Privacy Policy, letting Provider legally run its business and train non-AI algorithms without customer permission at every turn.

The Internal Ownership Trap.

Source · IP Assignment Deed & AI Employee Handbook · Extraction · §2.2 Conditional Assignment & The Thaler Bridge

"2.2 Conditional Assignment — AI-Generated Output. [...] To the extent that any Assigned IP is ultimately determined to be ineligible for copyright protection: (a) the assignment applies to the fullest extent legally permissible; (b) any rights that do exist — trade secret protection, contractual rights, sui generis rights — are assigned to the Company; [...] (d) if future legislation or judicial decisions establish new intellectual property protections for AI-assisted works, those rights are automatically assigned to the Company under this Deed without the need for a further instrument."

Intra-document architecture — THE "INSIDE" TRAP

Standard templates assign "all copyrights." If a court rules AI output cannot be copyrighted, standard assignments transfer nothing. This dependency chain mathematically secures enterprise IP regardless of how copyright law evolves.

NODE	VERBATIM TRIGGER	STRATEGIC ARCHITECTURE
§ 1.3 Assigned IP	"Assigned IP means... Human Creative Contributions — the Assignor's prompts, prompt sequences, prompt engineering decisions, selection logic, curation decisions, editorial modifications..."	LITIGATION DEFENCE Explodes standard IP definitions to explicitly capture prompt engineering and editorial curation as independent, assignable assets — the evidentiary basis to satisfy human-authorship requirements in <i>Thaler v. Perlmutter</i> .
§ 2.2 Conditional Assignment	"If future legislation or judicial decisions establish new intellectual property protections... those rights are automatically assigned to the Company."	THE THALER BRIDGE If AI-assisted work is ineligible for copyright, this conditional trap converts the assignment into trade-secret protection and proactively captures any future legislative rights — the enterprise never loses the asset's commercial value.
§ 1.6 Reversion Risk	"To prevent reversion, this Deed explicitly assigns all Assigned IP in perpetuity, for all territories worldwide, and for all modes of exploitation."	COMMERCIAL LEVERAGE Defeats the Indian Copyright Act §19(4) reversion trap. Undefined assignments revert to the creator after five years — this node forces precise perpetuity and multi-modal language to permanently protect portfolio valuation.

Ecosystem integration — THE "ACROSS" GATEWAYS

ORIGIN HOOK	TARGET CATCH (CONNECTED DOC)	ECOSYSTEM EFFECT
DOC_HND § 4.1 Policy Bridge	DOC_IP Recitals "The Company engages the Assignor to perform work that may involve the use of artificial intelligence tools... The Parties intend for this Deed to assign to the Company all intellectual property rights..."	The Policy-to-Legal Bridge. The Handbook creates operational awareness and mandate that all AI work belongs to the company; the Deed physically executes the transfer. In dispute, this proves the employee was fully informed before executing the assignment.
DOC_IP § 4.2 Evidence Bridge	DOC_SOP HITL Authorship Protocol — mandating version control, prompt logging, and draft preservation.	The Chain-of-Title Gateway. Assigning IP is useless if the company cannot prove a human created it. This gateway contractually forces employees to generate the exact evidentiary trail required to register a copyright and defend it against third-party challenges.

§ IV.01 — DIAGNOSTIC BACKBONE

The diagnostic backbone I use to isolate failure modes before drafting clauses.

Commercial contracts fail when practitioners draft for symptoms rather than system architecture. To bridge software engineering and legal liability, I have built an 80-point AI Threat Registry — classifying a deployment across three risk surfaces (Internal, External, Universal) and mapping it to ten technical archetypes (The Doer, The Judge, The Reader, etc.) to isolate exact failure modes before contractual firewalls are drafted.

The curated seven below map technical failure modes to multi-jurisdictional liability impact and the required contractual fix. Full registry available for in-room walkthrough.

THREAT INDEX · PROFILE	CLASSIFICATION	COMMERCIAL LIABILITY SYNTHESIS	ARCHITECTURAL FIX
I07_BIO_001 TIER · T1 EXTINCTION DEPTH · CORPORATE	The Diarization Voiceprint Trap BIPA · ILLINOIS	MECHANISM — Product performs speaker diarization on US state audio without BIPA §15 written consent. LAW — Strict liability under <i>Cruz v. Fireflies.AI</i> . No actual harm required. IMPACT — \$1,000–\$5,000 statutory damages per speaker per session. Near-automatic class certification.	DOC_AUP § 3.6 Pre-built BIPA §15-compliant speaker-consent module.
UNI_PRV_001 TIER · T2 UNCAPPED DEPTH · CORPORATE	Illegal Cross-Border Data Migration SCHREMS II · GDPR	MECHANISM — Processing EU user data on US servers without current SCCs and a Transfer Impact Assessment. LAW — <i>Schrems II</i> invalidation of Privacy Shield. IMPACT — GDPR penalties up to €20M or 4% of global revenue, plus orders suspending EU data flows.	DOC_DPA § 6.2 Pre-built Schrems II-compliant SCC and TIA module.
I03_HRM_001 TIER · T1 EXTINCTION DEPTH · CORPORATE	The Therapeutic Trap COMPANION AI · PRODUCT LIABILITY	MECHANISM — Product forms persistent emotional engagement without crisis-intervention protocols for self-harm signals. LAW — <i>Garcia v. Character.AI</i> establishes product-liability framing for companion AI. IMPACT — Wrongful-death damages and strict product-liability exposure, bypassing negligence defences.	DOC_AUP § 3.5 Pre-built Garcia-aware crisis-intervention module.

— § IV.01 · CONTINUED

Agent authority, fraud, and algorithmic collusion — the final three.

THREAT INDEX · PROFILE	CLASSIFICATION	COMMERCIAL LIABILITY SYNTHESIS	ARCHITECTURAL FIX
I02_DEC_001 TIER · T2 UNCAPPED DEPTH · CORPORATE	Vendor Immunity & HITL Theater CONSEQUENTIAL DECISIONS	MECHANISM — Product drives hiring/lending decisions with "human in the loop" that is functionally just clicking approve. LAW — <i>Mobley v. Workday</i> rejects AI vendor immunity, allowing suits under "agency" theory. IMPACT — Direct liability for discrimination damages; class action alongside the employer. "We just provide the tool" is dead.	DOC_AGT § 2.2 Pre-built Mobley-resilient decision-architecture module.
UNI_LIA_007 TIER · T2 UNCAPPED DEPTH · CORPORATE	Electronic Agent Authority UETA §14	MECHANISM — Automated system finalises transactions, approvals, or executions without a human reviewing the specific action. LAW — Uniform Electronic Transactions Act §14. IMPACT — Principal bound by agent's actions as a matter of law. Every transaction as if signed manually — liability compounds with scale.	DOC_TOS § 9.2 Pre-built UETA §14 principal-agent module.
UNI_FRD_001 TIER · T1 EXTINCTION DEPTH · CRIMINAL	AI Capability Misrepresentation SEC · "AI WASHING"	MECHANISM — Investor-facing representations about AI capabilities do not match operational reality (e.g., automation rate inflated). LAW — SEC designation of "AI washing" as civil and criminal fraud priority (<i>SEC v. Delphia</i> , <i>SEC+DOJ v. Nate Inc.</i>). IMPACT — SEC penalties, disgorgement, and direct criminal fraud exposure for named officers and founders.	DOC_AUP § 2.8 Pre-built Nate-aware capability-disclosure module.
I09_TRD_004 TIER · T2 UNCAPPED DEPTH · CRIMINAL	Algorithmic Collusion SHERMAN ACT §1	MECHANISM — AI pricing tool is used by multiple competitors and shares data, weights, or coordinated outputs — illegal price coordination. LAW — Sherman Act §1 and California AB 325. IMPACT — Treble damages and criminal antitrust exposure reaching vendor leadership, not just the customer.	DOC_AUP § 2.9 Pre-built Sherman-§1-resilient customer-siloed pricing module.

§ IV.02 — LIVE DEMONSTRATION

A black-box forensic run on a publicly available enterprise footprint.

DISCLAIMER — This is an independent, "black-box" diagnostic audit conducted exclusively on publicly available documentation (website, Terms of Service, Privacy Policy) as of April 2026. Sarvam AI is not a client. This is a demonstration of my diagnostic framework and does not constitute legal advice.

Standard legal due diligence reads contracts in a vacuum. My methodology maps a company's *Public Legal Posture* directly against its *Mechanical Reality* to find exactly where the liability architecture fractures. An automated forensic diagnostic script ingests a prospect's public footprint, maps their technical archetypes, and runs them deterministically against the 80-point Threat Registry. Running India's leading AI unicorn through this engine surfaced **21 structural gaps in under five minutes**.

PHASE 01 · TARGET INGESTION & PROFILE

SYSTEM LOG · JURISDICTIONAL MAPPING

```
// prospect ingestion
"prospect_meta": {
  "company_name": "Axonwise Private Limited",
  "hq_jurisdiction": "Bengaluru, India"
},
"jurisdictional_surface": [
  "EXT.03", "EXT.04",
  "EXT.08", "EXT.09",
  "EXT.10"
]
```

LIABILITY ASSESSMENT

The engine flagged Sarvam as a high-density, multi-jurisdictional target. Operating foundation models, B2B APIs, and B2C interfaces simultaneously automatically triggers compound regulatory exposure across Indian domestic law (DPDP, CPA 2019) and global extraterritorial frameworks (GDPR, CCPA) that generic SaaS contracts cannot cover.

PHASE 02 · CAPABILITY CLASSIFICATION

SYSTEM LOG · MECHANICAL CLASSIFICATION

```
<featuremap_scratchpad>
Feature: Samvaad Conversational Agents
Archetype: INT.01 (The Doer)
Reasoning: Acts autonomously on behalf
  of the user to fulfil customer intent.
Evidence: "Agents connect to your
  enterprise tools, use your data,
  take action."
</featuremap_scratchpad>

<featuremap_scratchpad>
Feature: Studio AI Dubbing & Translation
Archetype: INT.04 (The Creator)
Reasoning: Generates synthetic,
  localised audio and text outputs.
</featuremap_scratchpad>
```

LIABILITY ASSESSMENT

The engine strips away marketing adjectives and classifies products purely by their mechanical reality. A chatbot is no longer passive software — if it takes independent action, it is legally an Electronic Agent. This mapping dictates exactly which legal firewalls must be present in the Terms of Service.

— § IV.03 · PHASE 03 — THREAT MAPPING & REMEDIATION

Three fractures of highest immediate commercial and regulatory exposure.

From the 21 structural gaps surfaced by the diagnostic engine.

TEARDOWN 01	UNI_LIA_007 · Samvaad · Agentic Liability Fracture	PAIN TIER · T2
MECHANICAL TRIGGER — Sarvam markets Samvaad as an autonomous agent capable of executing enterprise actions (INT.01 "Doer"), but governs it with legacy Web 2.0 ToS lacking principal-agent authority boundaries or pre-commit reversal protocols. Remediation · Deploy Agentic Addendum and Output-Authority limitation (Work Sample 01).	LEGAL FRACTURE [IN] Uncapped exposure for unauthorised automated transactions under Contract Act 1872 and IT Act §10A. [US] Fails to bind end-user under UETA §14, leaving vendor strictly liable for third-party damages. [EU] Fails to establish human-oversight guardrails mandated for autonomous execution under the EU AI Act.	
TEARDOWN 02	UNI_HAL_001 · Samvaad · Hallucination Reliance Gap	PAIN TIER · T2
MECHANICAL TRIGGER — The AI surfaces generated commitments to users without a contemporaneous, conspicuous disclaimer at the point of interaction that outputs are probabilistic and non-binding. Remediation · Deploy Hallucination Waiver & Conditional Defence (Work Sample 01).	LEGAL FRACTURE [IN] Triggers actionable "deficiency in service" under CPA 2019 §2(11). [US] Negates defence against negligent misrepresentation (<i>Moffatt</i>), exposing direct reliance damages. [EU] Exposes deployer to strict liability under the revised Product Liability Directive for defective AI output.	
TEARDOWN 03	I05_PRV_002 · Akshar · Enterprise Training Data Gap	PAIN TIER · T1
MECHANICAL TRIGGER — Akshar digitises unstructured enterprise documents. The Privacy Policy relies on generic Web 2.0 data collection clauses and lacks explicit training-data opt-outs — implicitly allowing enterprise inputs to be baked into model weights. Remediation · Deploy unified DPA featuring the RAG-Only Mandate (Work Sample 02).	LEGAL FRACTURE [IN] B2B buyers in structural violation of DPDP erasure obligations — weight-encoded data cannot be selectively deleted ("algorithmic disgorgement"). [US] FTC Section 5 exposure for deceptive data collection, resulting in forced model deletion. [EU] Violates "Right to Erasure" under GDPR Article 17.	

The instruments that actually drive commercial AI practice.

Not a compliance encyclopedia — the working frame I use to track the live regulatory surface, surface interactions that generic frameworks miss, and map the active liability frontier for enterprise deployment.

INDIA · DOMESTIC BASELINE & FRONTIER		
INSTRUMENT	STATUS / TIMELINE	PRACTITIONER NOTE ON AI RELEVANCE
IT Act 2000 §43A & DPDP Act 2023	Phased · §43A repealed 13 May 2027	Post-May 2027, DPDP activates with penalties up to ₹250 Cr but no private right of damages. Shifts AI data exposure from private litigation to state regulatory risk, mandating strict RAG architecture to prevent immutable model-weight contamination.
Consumer Protection Act 2019	In force · 2-yr limitation	"Deficiency in service" (§2(11)) is the most direct statutory route for Indian consumers harmed by AI hallucinations. With a two-year limitation, AI startups are actively accruing unseen liabilities legacy SaaS contracts do not defend against.
Contract Act 1872 & Copyright Act 1957	In force	Contract Act governs enforceability of AI hallucination waivers. Copyright Act §19(4) — five-year assignment reversion — creates a structural trap for AI-assisted work product ownership if not explicitly contracted around in perpetuity.
MeitY AI Advisories & Draft IT Rules	Active consultation · Non-binding	March 2024 advisory sets articulated regulatory expectation for labelling the "inherent fallibility" of AI outputs. Ignoring it establishes a de facto baseline for negligence claims even before formal codification.
EUROPEAN UNION · EXTRATERRITORIAL OVERLAY		
INSTRUMENT	STATUS / TIMELINE	PRACTITIONER NOTE ON AI RELEVANCE
EU AI Act	Phased · Art. 50 from 2 Aug 2026	Risk-tiered architecture with global extraterritorial reach. Triggers immediate compliance obligations for Indian companies placing AI systems or APIs on the EU market.
GDPR (Reg. EU 2016/679)	In force	The core architectural friction for foundation models: Art. 17 (erasure) vs. immutable model weights; Art. 22 (automated decision-making); analytics routed via Independent Controller carve-outs to protect vendor revenue.
EU Product Liability Directive	Transposition by 9 Dec 2026	Reclassifies software and AI systems as "products." Establishes a strict liability pathway for plaintiffs suffering damage from defective AI outputs, bypassing traditional negligence hurdles.

— § V.01 · CONTINUED

The American commercial infrastructure and state-level fragmentation.

The US layer is not a single regime but a patchwork of commercial code, state privacy statutes, and employment law that every Indian AI vendor targeting US enterprise buyers must map clause-by-clause.

UNITED STATES · COMMERCIAL INFRASTRUCTURE & STATE-LEVEL FRAGMENTATION

INSTRUMENT	STATUS / TIMELINE	PRACTITIONER NOTE ON AI RELEVANCE
UCC Art. 2 & UETA §14	In force	UCC §2-316 dictates the conspicuousness standard for AS IS disclaimers. UETA §14 is foundational for agentic AI — binds a principal to contracts formed by their electronic agent without human review, requiring strict contractual permission scoping.
State Privacy Overlay (CCPA, VA, TX, CO)	In force · Rolling	Creates strict "service provider" obligations for AI deployments. Forces B2B Indian startups to architect contractual data boundaries between customer inputs and upstream model training.
EEOC Title VII & Local AI Legislation	In force · Rolling	Governs algorithmic discrimination (NYC LL 144 bias audits; Colorado AI Act). Imposes disparate-impact liability for AI-assisted employment decisions regardless of developer intent.

§ V.02 — APPLIED CASE LAW

The doctrinal anchors that hold the architecture to the ground.

Each of the five cases below anchors a specific firewall in the stack — its ruling maps to a named clause, not an abstract concern.

Moffatt v. Air Canada

Canada · CRT · 2024

NEGLIGENT MISREP

The anchor for negligent misrepresentation via AI. Kills the "separate tool" defence and establishes that an AI chatbot operates with representational authority for the deploying company.

Walters v. OpenAI

US · Ga. Super. · 2025

DISCLAIMER WEIGHT

Validates the evidentiary weight of specific, point-of-interaction AI disclaimers to defeat claims of "reasonable reliance" on hallucinated outputs.

Specht & Meyer

US · 2d Cir. · 2002 / 2017

CLICKWRAP
STANDARD

Mandates that AI output waivers must be presented with unambiguous affirmative assent, rendering buried "browsewrap" SaaS terms useless.

Thaler v. Perlmutter

US · D.D.C. · 2023

HUMAN AUTHORSHIP

Denies copyright protection to works generated autonomously by AI, mandating strict internal governance (HITL SOPs) to prove human creative contribution.

Bartz v. Anthropic & NYT v. OpenAI

US · N.D. Cal. / SDNY

TRAINING-DATA
RECKONING

The training-data reckoning. Bartz (\$1.5B settlement) signals erosion of the fair-use defence, forcing downstream B2B wrappers to demand robust upstream IP indemnities.

The doctrine is the anchor. The clause is the line. The ecosystem is the ship.

Probabilistic technology has outpaced deterministic legal frameworks.

My foundation in litigation trained me to navigate uncodified risk, but my value to this firm is defined entirely by my present execution.

Enterprises are currently absorbing massive, unquantified liability because probabilistic technology has fundamentally outpaced deterministic legal frameworks.

I am not burdened by legacy corporate templates, because the legal architecture required to govern generative AI did not exist two years ago. It had to be built from the ground up. I did not wait for the market to establish a precedent, nor did I wait to be formally trained on the solution.

Over the last several months, I executed a rigorous, independent deep-dive into the intersection of AI architecture, the incoming DPDP regime, and global enforcement trends to engineer the frameworks presented in this dossier.

This portfolio is the direct proof of that capability. It bridges the gap between raw technological risk and commercial contract defensibility across output liability, data governance, and IP preservation. I built this infrastructure preemptively because my intent is to deliver senior-level commercial value to a firm on day one.

Now, I am ready to deploy it.

CORRESPONDENCE

shwetav.alpna.singh
@gmail.com

THE FULL ECOSYSTEM

Ten production-ready architectures, an 80-point threat registry, and the reasoning chains between them — available for in-room walkthrough.

COLOPHON

Shwetav Singh · Advocate
Portfolio & Proof of Competence
MMXXVI · Vol. 01